

# An Introduction into Parallelization

## Part II - Design

Jascha Schewtschenko

Royal Observatory of Edinburgh, University of Edinburgh

May 16, 2025



# Outline

## 1 Parallel Programming Models (cont.)

- Recap: Computer Architectures
- Distributed Parallelism with Message Passing
- Hybrid Models

## 2 Designing Parallel Programs

- Understand your problem and tools
- Partitioning - Domain vs functional decomposition
- Data Dependence / Race conditions
- Synchronization
- Communication
- Load balancing
- I/O
- Performance Analysis & Tuning
- Conclusions / tl;dr



# Parallel Programming Models



## THREE LEVELS OF PARALLEL PROGRAMMING



**VECTORIZATION**



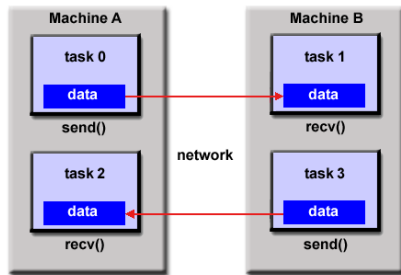
**MULTITHREADING**



**DISTRIBUTED  
PARALELLISM**

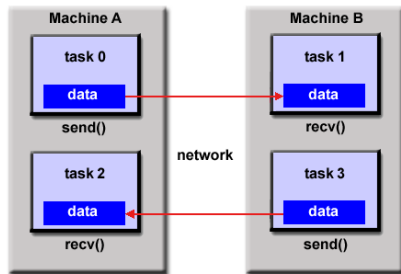
# Distributed Parallelism with Message Passing

- set of processes with own local memory - reside on the same node and/or across multiple nodes.



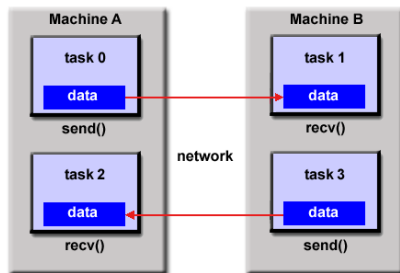
# Distributed Parallelism with Message Passing

- set of processes with own local memory - reside on the same node and/or across multiple nodes.
- Data between processes is exchanged through sending and receiving messages (“Message Passing”)




# Distributed Parallelism with Message Passing

- set of processes with own local memory - reside on the same node and/or across multiple nodes.
- Data between processes is exchanged through sending and receiving messages (“Message Passing”)
- The message passing requires cooperation between processes (each send must have a matching receive operation)



# Distributed Parallelism with Message Passing

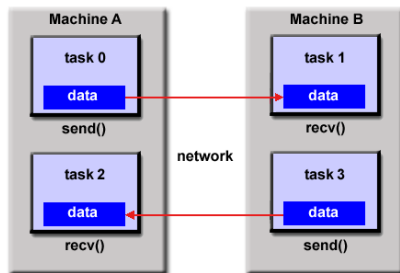
- set of processes with own local memory - reside on the same node and/or across multiple nodes.
- Data between processes is exchanged through sending and receiving messages (“Message Passing”)
- The message passing requires cooperation between processes (each send must have a matching receive operation)
- The message system can also be used to synchronize processes






# Distributed Parallelism with Message Passing

- set of processes with own local memory - reside on the same node and/or across multiple nodes.
- Data between processes is exchanged through sending and receiving messages (“Message Passing”)
- The message passing requires cooperation between processes (each send must have a matching receive operation)
- The message system can also be used to synchronize processes
- *MPI* is the “de facto” standard



# Hybrid Models

- Allows to make best use of locally shared memory or hardware, while still allowing for a good scalability across multiple nodes
- Comes with a significant increase in complexity/costs
- certain incompatibilities between libraries may exist (e.g. lack of thread-safety of MPI library)




# Designing Parallel Programs



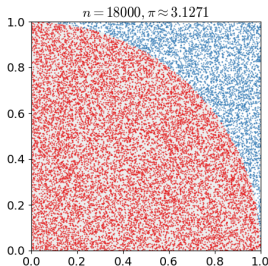
# Understand your problem (!)

- some problems can be easily parallelized e.g. calculation of  $\pi$  by Monte-Carlo sampling



# Understand your problem (!)

- some problems can be easily parallelized e.g. calculation of  $\pi$  by Monte-Carlo sampling

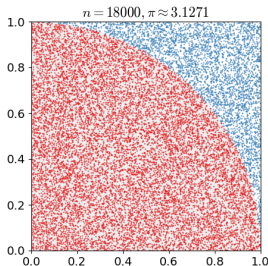


- others allow for little-to-no parallelism e.g. recursive calculation of Fibonacci series

$$F(n) = F(n-1) + F(n-2), \quad F(0) = 0, \quad F(1) = 1$$

# Understand your problem (!)

- some problems can be easily parallelized e.g. calculation of  $\pi$  by Monte-Carlo sampling



- others allow for little-to-no parallelism e.g. recursive calculation of Fibonacci series

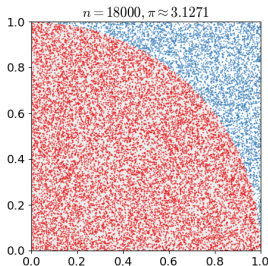
$$F(n) = F(n-1) + F(n-2), \quad F(0) = 0, \quad F(1) = 1$$

- Identify inhibitors to parallelism: data-dependencies, I/O bottlenecks



# Understand your problem (!)

- some problems can be easily parallelized e.g. calculation of  $\pi$  by Monte-Carlo sampling



- others allow for little-to-no parallelism e.g. recursive calculation of Fibonacci series

$$F(n) = F(n-1) + F(n-2), \quad F(0) = 0, \quad F(1) = 1$$

- Identify inhibitors to parallelism: data-dependencies, I/O bottlenecks
- Consider replacing your algorithms with equivalent ones better suited for parallelism



# Understand your tools/program (!)

- pick optimal parallel programming model for your infrastructure





# Understand your tools/program (!)

- pick optimal parallel programming model for your infrastructure
- make use of hardware optimization (e.g. vectorization, optimized libraries like MKL)



# Understand your tools/program (!)

- pick optimal parallel programming model for your infrastructure
- make use of hardware optimization (e.g. vectorization, optimized libraries like MKL)
- identify hotspots in your program, i.e. routines where program spends lots of time in and check for improvement in parallelism ( $\rightarrow$  Amdahl's Law/Scaling)



# Partitioning - Domain decomposition

- In this type of partitioning, the data associated with a problem is decomposed. Each parallel task then works on a portion of the data.



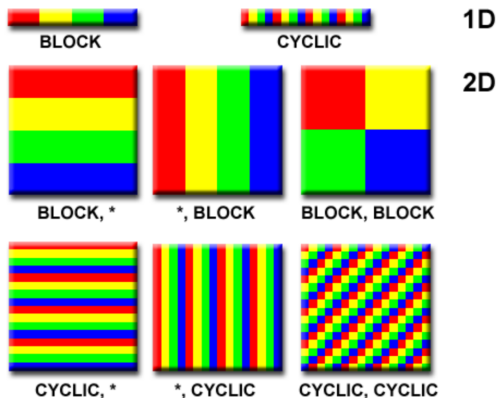
# Partitioning - Domain decomposition

- In this type of partitioning, the data associated with a problem is decomposed. Each parallel task then works on a portion of the data.
- There are different ways to partition data:




# Partitioning - Domain decomposition

- In this type of partitioning, the data associated with a problem is decomposed. Each parallel task then works on a portion of the data.
- There are different ways to partition data:



# Partitioning - Domain decomposition

- In this type of partitioning, the data associated with a problem is decomposed. Each parallel task then works on a portion of the data.
- There are different ways to partition data:




# Partitioning - functional decomposition

- In this approach, the focus is on the computation that is to be performed rather than on the data manipulated by the computation.



# Partitioning - functional decomposition

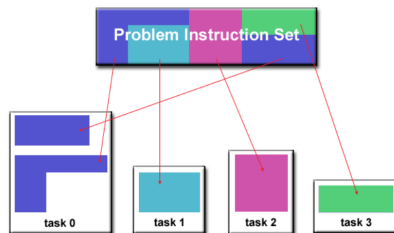
- In this approach, the focus is on the computation that is to be performed rather than on the data manipulated by the computation.
- Tasks/threads then specialize in doing specific parts of the overall work:





# Partitioning - functional decomposition

- In this approach, the focus is on the computation that is to be performed rather than on the data manipulated by the computation.
- Tasks/threads then specialize in doing specific parts of the overall work:



- implemented e.g. in master/slave paradigm (see exercises)

# Data Dependence

- *data dependence* results from multiple use of the same memory location by different tasks.



# Data Dependence

- *data dependence* results from multiple use of the same memory location by different tasks.
- Dependencies are important to parallel programming because they are one of the primary inhibitors to parallelism (cf. Fibonacci series).



# Data Dependence

- *data dependence* results from multiple use of the same memory location by different tasks.
- Dependencies are important to parallel programming because they are one of the primary inhibitors to parallelism (cf. Fibonacci series).
- This can also cause a so called race condition:

Thread 1	Thread 2		Integer value
			0
read value		←	0
increase value			0
write back		→	1
	read value	←	1
	increase value		1
	write back	→	2

Thread 1	Thread 2		Integer value
			0
read value		←	0
	read value	←	0
increase value			0
	increase value		0
write back		→	1
	write back	→	1

# Synchronization

- synchronization is used to control the flow of processes/threads to deal with “mutually exclusive” resources (using locks/semaphores) to resolve e.g. race conditions



# Synchronization

- synchronization is used to control the flow of processes/threads to deal with “mutually exclusive” resources (using locks/semaphores) to resolve e.g. race conditions
- ... or to synchronize the calculations of processes (using barriers) for communication to exchange results or to redistribute the workload



# Communication

- Depending on the type of problem and used decomposition approach, more or less communication may be required



# Communication

- Depending on the type of problem and used decomposition approach, more or less communication may be required
- e.g. for *embarrassingly parallel* problem tasks run (mostly) independent of each other while e.g. in cosmological simulations long-range forces have to be communicated between tasks working on neighbouring domains and particles may have to be exchanged if domain borders or particles move.





# Communication

- Depending on the type of problem and used decomposition approach, more or less communication may be required
- e.g. for *embarrassingly parallel* problem tasks run (mostly) independent of each other while e.g. in cosmological simulations long-range forces have to be communicated between tasks working on neighbouring domains and particles may have to be exchanged if domain borders or particles move.
- Factors to consider:



# Communication

- Depending on the type of problem and used decomposition approach, more or less communication may be required
- e.g. for *embarrassingly parallel* problem tasks run (mostly) independent of each other while e.g. in cosmological simulations long-range forces have to be communicated between tasks working on neighbouring domains and particles may have to be exchanged if domain borders or particles move.
- Factors to consider:
  - ▶ Communication overhead



# Communication

- Depending on the type of problem and used decomposition approach, more or less communication may be required
- e.g. for *embarrassingly parallel* problem tasks run (mostly) independent of each other while e.g. in cosmological simulations long-range forces have to be communicated between tasks working on neighbouring domains and particles may have to be exchanged if domain borders or particles move.
- Factors to consider:
  - ▶ Communication overhead
  - ▶ Latency vs. Bandwidth



# Communication

- Depending on the type of problem and used decomposition approach, more or less communication may be required
- e.g. for *embarrassingly parallel* problem tasks run (mostly) independent of each other while e.g. in cosmological simulations long-range forces have to be communicated between tasks working on neighbouring domains and particles may have to be exchanged if domain borders or particles move.
- Factors to consider:
  - ▶ Communication overhead
  - ▶ Latency vs. Bandwidth
  - ▶ Synchronous vs. asynchronous communications




# Communication

- Depending on the type of problem and used decomposition approach, more or less communication may be required
- e.g. for *embarrassingly parallel* problem tasks run (mostly) independent of each other while e.g. in cosmological simulations long-range forces have to be communicated between tasks working on neighbouring domains and particles may have to be exchanged if domain borders or particles move.
- Factors to consider:
  - ▶ Communication overhead
  - ▶ Latency vs. Bandwidth
  - ▶ Synchronous vs. asynchronous communications
  - ▶ Point-to-Point vs. collective communications



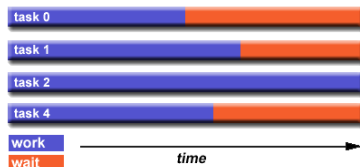
# Load balancing

- Load balancing needed to ensure that processors are optimally i.e. minimizing idle times at synchronization points



# Load balancing


- Load balancing needed to ensure that processors are optimally i.e. minimizing idle times at synchronization points



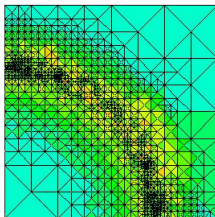
- requires well-balanced workload distribution between processors

# Load balancing

- Load balancing needed to ensure that processors are optimally i.e. minimizing idle times at synchronization points



- requires well-balanced workload distribution between processors
- difficult in heterogeneous, dynamic problem sets with incomplete information about the actual workload






## Load balancing (cont.)

- requires repeated repartitioning the problem based on estimate of workload



# Load balancing (cont.)

- requires repeated repartitioning the problem based on estimate of workload
- alternatively, use asynchronous approach with scheduler-task pool with smaller workload packages



- data transfer to storage medias, network file servers

- data transfer to storage medias, network file servers
- I/O operations are generally obstacles for parallelism as they require orders of magnitude more time than memory operations.

- data transfer to storage medias, network file servers
- I/O operations are generally obstacles for parallelism as they require orders of magnitude more time than memory operations.
- Like for shared memory, may require synchronization (e.g. to avoid two processes (over)writing the same file simultaneously in a shared file system

- data transfer to storage medias, network file servers
- I/O operations are generally obstacles for parallelism as they require orders of magnitude more time than memory operations.
- Like for shared memory, may require synchronization (e.g. to avoid two processes (over)writing the same file simultaneously in a shared file system
- Parallel filesystems available (e.g. GPFS, Lustre, HDFS)



- data transfer to storage medias, network file servers
- I/O operations are generally obstacles for parallelism as they require orders of magnitude more time than memory operations.
- Like for shared memory, may require synchronization (e.g. to avoid two processes (over)writing the same file simultaneously in a shared file system
- Parallel filesystems available (e.g. GPFS, Lustre, HDFS)
- more on this in the “HPC & (Big) Data” talk!



# Performance Analysis & Tuning

- Analysing and tuning parallel program performance can be much more challenging than for serial programs as interactions between tasks result in very complex dynamics





# Performance Analysis & Tuning

- Analysing and tuning parallel program performance can be much more challenging than for serial programs as interactions between tasks result in very complex dynamics
- There are a number of excellent tools for this task: e.g. Intel VTune Amplifier and Intel Trace Analyzer, ARM DDT



# Performance Analysis & Tuning

- Analysing and tuning parallel program performance can be much more challenging than for serial programs as interactions between tasks result in very complex dynamics
- There are a number of excellent tools for this task: e.g. Intel VTune Amplifier and Intel Trace Analyzer, ARM DDT
- Unfortunately, covering this topic in all the detail would go beyond the scope of this introduction to parallel program, so we focus on a simple python profiler as an example, Yappi.



# Performance Analysis & Tuning

- Analysing and tuning parallel program performance can be much more challenging than for serial programs as interactions between tasks result in very complex dynamics
- There are a number of excellent tools for this task: e.g. Intel VTune Amplifier and Intel Trace Analyzer, ARM DDT
- Unfortunately, covering this topic in all the detail would go beyond the scope of this introduction to parallel program, so we focus on a simple python profiler as an example, Yappi.

Two caveats:



# Performance Analysis & Tuning

- Analysing and tuning parallel program performance can be much more challenging than for serial programs as interactions between tasks result in very complex dynamics
- There are a number of excellent tools for this task: e.g. Intel VTune Amplifier and Intel Trace Analyzer, ARM DDT
- Unfortunately, covering this topic in all the detail would go beyond the scope of this introduction to parallel program, so we focus on a simple python profiler as an example, Yappi.

Two caveats:

- **Heisenberg's uncertainty principle**

*The more precisely you know the position of something, the less precisely you can tell its speed and momentum.*

applies to profiling, too!



# Performance Analysis & Tuning

- Analysing and tuning parallel program performance can be much more challenging than for serial programs as interactions between tasks result in very complex dynamics
- There are a number of excellent tools for this task: e.g. Intel VTune Amplifier and Intel Trace Analyzer, ARM DDT
- Unfortunately, covering this topic in all the detail would go beyond the scope of this introduction to parallel program, so we focus on a simple python profiler as an example, Yappi.

Two caveats:

- **Heisenberg's uncertainty principle**

*The more precisely you know the position of something, the less precisely you can tell its speed and momentum.*

applies to profiling, too!

- profilers do **NOT** tell you **HOW** to optimize your code, they merely tell you **WHERE** to (potentially) do so for the best yield.



# Performance Analysis & Tuning - Yappi

- $\mathbf{Y}_{(et)}\mathbf{A}_{(nother)}\mathbf{P}_{(ython)}\mathbf{P}_{(rof)}\mathbf{I}_{(ler)}$  supports multi-threaded python out of the box; allows for easy CPU-time measurements



# Performance Analysis & Tuning - Yappi

- $Y_{(et)}A_{(nother)}P_{(ython)}P_{(rof)}I_{(ler)}$  supports multi-threaded python out of the box; allows for easy CPU-time measurements
- (multi)processes have to be profiled seperately (cf. Exercises)



# Performance Analysis & Tuning - Yappi

- $Y_{(et)}A_{(nother)}P_{(ython)}P_{(rof)}I_{(ler)}$  supports multi-threaded python out of the box; allows for easy CPU-time measurements
- (multi)processes have to be profiled separately (cf. Exercises)
- exports results into callgrind & pstat format (to be loaded by analyzers/visualisation tools)

```
Function stats for (.MainThread) (0)
Clock type: WALL
Ordered by: totaltime, desc

name                                ncall      tsub      ttot      tavg
..hon3.8/threading.py:859 Thread.run 2      0.000017  44.09970  22.04985
..ads_revisited/helper.py:26 runall  2      0.002986  44.09969  22.04984
time_trials_threading.py:11 worker  10     0.052127  44.09676  4.409676
..ads_revisited/helper.py:10 foreach 1      0.000022  22.06303  22.06303
..on3.8/threading.py:979 Thread.join 2      0.000016  22.06242  11.03121
...1017 Thread_wait_for_tstate_lock 2      0.000014  22.06240  11.03120
time_trials_threading.py:23 g        10     0.084264  9.152582  0.915258
..time_trials_threading.py:24 _listcomp 10     7.603336  9.068318  0.906832
time_trials_threading.py:26 h        10     7.371068  8.857567  0.885757
time_trials_threading.py:19 f        10     5.973178  5.973178  0.597318
..ges/numpy/core/numeric.py:144 ones 10     0.000200  0.045074  0.004507
..ray.function _internals>:2 copyto 10     0.000135  0.044414  0.004411
..n3.8/threading.py:834 Thread.start 2      0.000023  0.000501  0.000250
..hon3.8/threading.py:540 Event.wait 2      0.000014  0.000375  0.000188
..8/threading.py:270 Condition.wait 2      0.000027  0.000350  0.000175
..ads_revisited/helper.py:52 _listcomp 1      0.000011  0.000091  0.000091
..8/threading.py:761 Thread._init_ 2      0.000039  0.000080  0.000040
..8/threading.py:505 Event._init_ 2      0.000006  0.000021  0.000010
..mpy/core/multiarray.py:1043 copyto 10     0.000018  0.000018  0.000002
..n3.8/threading.py:944 Thread._stop 2      0.000014  0.000016  0.000008
..hreading.py:222 Condition._init_ 2      0.000014  0.000014  0.000007
..reading.py:261 Condition._is_owned 2      0.000005  0.000009  0.000005
..8/threading.py:1306 _current_thread 4      0.000008  0.000009  0.000002
..n3.8/_weakrefset.py:81 WeakSet.add 2      0.000007  0.000008  0.000004
..reading.py:246 Condition._enter_ 2      0.000004  0.000006  0.000003
..ing.py:255 Condition._release_save 2      0.000004  0.000005  0.000002
..hreading.py:249 Condition._exit_ 2      0.000005  0.000005  0.000002
..ython3.8/_weakrefset.py:38 _remove 2      0.000004  0.000004  0.000002
..py:258 Condition._acquire_restore 2      0.000001  0.000003  0.000002
..ython3.8/threading.py:734 _newname 2      0.000003  0.000003  0.000002
..np.py:1177 _make_invoke_excpthook 2      0.000003  0.000003  0.000002
..reading.py:1095 _MainThread.daemon 4      0.000002  0.000002  0.000000
..n3.8/threading.py:513 Event.is_set 4      0.000001  0.000001  0.000000
```





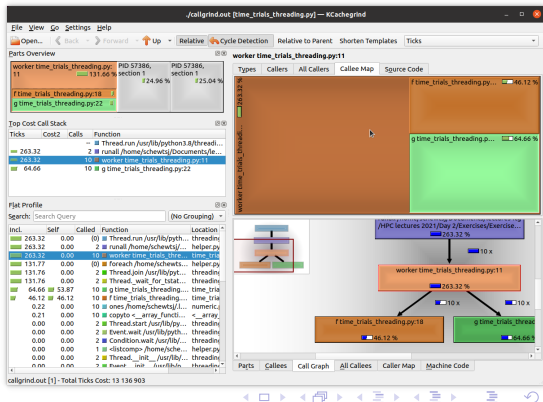
# Performance Analysis & Tuning - Yappi

- $Y_{(et)} A_{(nother)} P_{(ython)} P_{(rof)} I_{(ler)}$  supports multi-threaded python out of the box; allows for easy CPU-time measurements
- (multi)processes have to be profiled separately (cf. Exercises)
- exports results into callgrind & pstat format (to be loaded by analyzers/visualisation tools)

Function stats for (.MainThread) (0)

Clock type: WALL  
Ordered by: totaltime, desc

name	ncall	tsub	ttot	tavg
..hon3.8/threading.py:859 Thread.run	2	0.000617	44.09970	22.04985
..eds_revisited/helper.py:26 runall	2	0.002986	44.89969	22.04984
time_trials_threading.py:11 worker	10	0.052127	44.09676	4.409676
..ads_revisited/helper.py:18 foreach	1	0.000622	22.06303	22.06303
..on3.8/threading.py:979 Thread.join	2	0.000616	22.06242	11.03121
..1017 Thread_wait_for_tstate_lock	2	0.000614	22.06240	11.03120
time_trials_threading.py:23 g	10	0.084264	9.152582	0.915258
..trials_threading.py:24 <listcomp>	10	7.603336	9.068318	0.906832
time_trials_threading.py:26 h	7	7.371018	8.85657	0.856567
time_trials_threading.py:19 f	5	5.973178	5.973178	0.597318
..ges/numpy/core/numeric.py:144 ones	10	0.000200	0.045074	0.004507
..ray_function_intervals>:2 copyto	10	0.000135	0.044414	0.004411
..n3.8/threading.py:834 Thread.start	2	0.000023	0.000501	0.000250
..hon3.8/threading.py:540 Event.wait	2	0.000614	0.000375	0.000188
..8/threading.py:270 Condition.wait	2	0.000027	0.000350	0.000175
..revisited/helper.py:52 <listcomp>	1	0.000011	0.000091	0.000091
..8/threading.py:761 Thread._init	2	0.000039	0.000080	0.000040
..8/threading.py:505 Event._init	2	0.000006	0.000021	0.000010
..npy/core/multibyte.py:1043 copyto	10	0.000018	0.000018	0.000002
..n3.8/threading.py:944 Thread._stop	2	0.000014	0.000016	0.000008
..threading.py:222 Condition._init	2	0.000014	0.000014	0.000007
..reading.py:261 Condition._is_owned	2	0.000005	0.000009	0.000005
..8/threading.py:1306 current_thread	4	0.000008	0.000009	0.000002
..n3.8/weakrefset.py:181 WeakSet.add	2	0.000007	0.000003	0.000004
..reading.py:246 Condition._enter	2	0.000004	0.000006	0.000003
..ing.py:255 Condition._release_save	2	0.000004	0.000005	0.000002
..threading.py:249 Condition._exit	2	0.000005	0.000005	0.000002
..thon3.8/weakrefset.py:38 _remove	2	0.000004	0.000004	0.000002
..py:258 Condition._acquire_restore	2	0.000001	0.000003	0.000002
..thon3.8/threading.py:734 _newname	2	0.000003	0.000003	0.000002
..np.py:1177 make_invoke_excepthook	2	0.000003	0.000003	0.000002
..reading.py:1095 _MainThread.daemon	4	0.000002	0.000002	0.000000
..n3.8/threading.py:513 Event._is_set	4	0.000001	0.000001	0.000000



# Conclusions / tl;dr

- Writing code that does the job is one thing; writing it in an optimized way is a different game altogether  $\Rightarrow$  HPC/HTC requires highly optimized code for “best” performance!



# Conclusions / tl;dr

- Writing code that does the job is one thing; writing it in an optimized way is a different game altogether  $\Rightarrow$  HPC/HTC requires highly optimized code for “best” performance!
- Parallelization techniques provide ways to reduce the “walltime” of a code (usually at the cost of an increased “CPU time”); requires in-depth knowledge in parallel algorithms



# Conclusions / tl;dr

- Writing code that does the job is one thing; writing it in an optimized way is a different game altogether  $\Rightarrow$  HPC/HTC requires highly optimized code for “best” performance!
- Parallelization techniques provide ways to reduce the “walltime” of a code (usually at the cost of an increased “CPU time”); requires in-depth knowledge in parallel algorithms
- often, there are already “low-hanging fruits” in the serial code (reduction of I/O, amount of function calls, object conversions, etc.); demands in-depth knowledge of how things work “under the hood”.



# Conclusions / tl;dr

- Writing code that does the job is one thing; writing it in an optimized way is a different game altogether  $\Rightarrow$  HPC/HTC requires highly optimized code for “best” performance!
- Parallelization techniques provide ways to reduce the “walltime” of a code (usually at the cost of an increased “CPU time”); requires in-depth knowledge in parallel algorithms
- often, there are already “low-hanging fruits” in the serial code (reduction of I/O, amount of function calls, object conversions, etc.); demands in-depth knowledge of how things work “under the hood”.
- tools like profilers can help you to identify what to optimize, but you still need the knowledge on how to do this.

